



# Telling the story of WeissBeerger Cloud Evolution

Eitan Sela - System Architect

WeissBeerger

A part of the **ABInBev** Family

[eitan.sela@weissbeerger.com](mailto:eitan.sela@weissbeerger.com)

WeissBeerger

BEVERAGE  
ANALYTICS

# \$ whoami

- "Hands-On" system Architect with more than 18 years of experience with billing, banking, information security (DLP) and Cloud IoT/Big Data applications.
- Big Data specialist – Hadoop, Spark, Hive and EMR on AWS.
- Work with vast AWS services, and with serverless projects especially.
- Java development, scalability performance and stabilization expert.
- Alexa skills developer.
- Love to share my experience in lectures and meetups.



Israel Football

by Eitan Sela

★★★★★ 1

Free to Enable

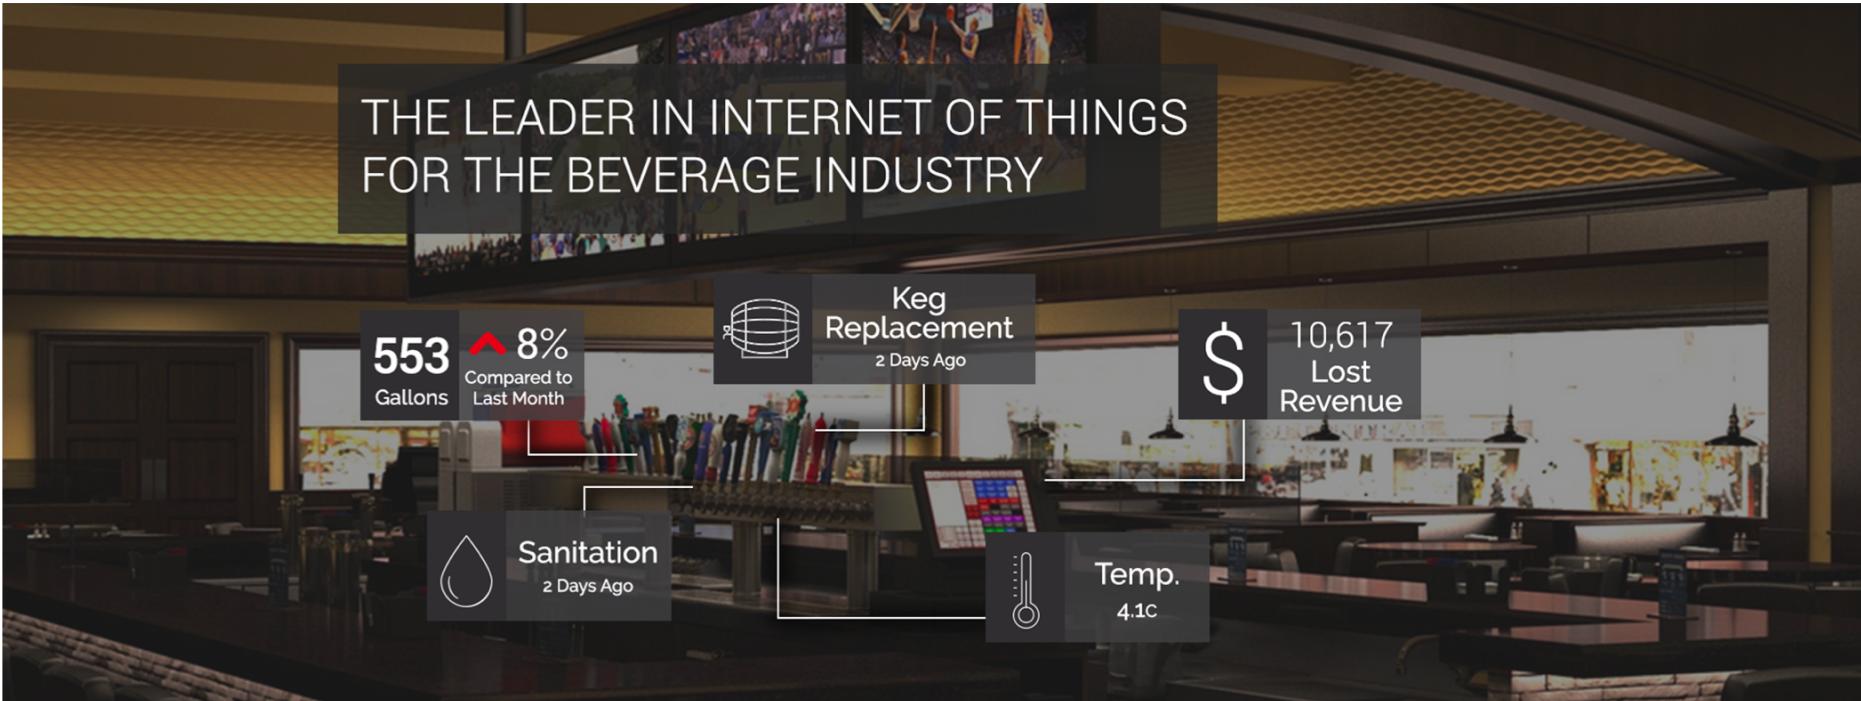
"Alexa, open Israel Football"

## What to expect from this session

- WeissBeriger – Company overview & use case.
- Architecture - August 2017.
- What were the problems we faced back then?
- Cloud Evolution – Building the puzzle.
- Architecture - December 2018.
- Future Challenges.

# WeissBeerger use case

- WeissBeerger bridges the gap between breweries, bars and customers.



# How does it work?

- IoT (Pouring) – Beverage Analytics Hub.
- Point of sales – POS Vendors, via REST API, S3, DB, etc.



# WeissBeerger Architecture - August 2017

# Web Dashboard and Mobile App Architecture – August 2017

Monolith PHP



Amazon EC2



# But it is actually a huge monolith EC2! And all Apps are stateful!

**Monolith PHP**  
Dashboard and  
Mobile App  
(stateful)

**Monolith Java**  
Point of sale  
App (stateful)



**Data science Algo**  
(stateful)

**MongoDB**  
For Data science  
Model

# AWS Usage – August 2017

- EC2
- Route 53



# **What were the problems we faced back then?**

## Problems with the architecture

- One monolithic EC2 – **No high availability & redundancy, manual security updates.**
- Apps are Monolithic and Stateful – **No reuse & modularity, Can not be scaled!**
  - IoT HTTP Gateway - Developed in C++, stateful and un-maintainable.
  - POS – Developed in pure Java, with custom thread handling, in Spring Boot.
- Apps are updated manually in the server with git pull – **No automation - CI/CD.**

## Problems with the architecture – con't

- Security:
  - Development and Production EC2 are on the same account.
  - No dedicated security groups and vpc's.
  - Only one subnet open to the internet.
- No Monitoring:
  - Developers are reactive (vs. proactive).
  - No one knows if something fails in production.

The main problem: Poor performance of Dashboard and Mobile App



# Cloud Evolution – Building the puzzle

# Re-Architecture project – June 2017

Principle	Concepts
Security	<ul style="list-style-type: none"><li>• Network &amp; Infra Security</li><li>• Secure software implementation practices</li><li>• Penetration testing</li><li>• Applicative security using roles and permissions model</li></ul>
Robust, Scalable & Performant	<ul style="list-style-type: none"><li>• Ability to scale out to meet demand</li><li>• APIs, Services &amp; Applications built to perform on Day 1</li></ul>
High Availability & Redundancy	<ul style="list-style-type: none"><li>• Multiple Servers in Multiple Physical Locations</li><li>• Reduce single points of failure</li><li>• Automated backup and restore</li><li>• Designed for fast failover</li></ul>
Be Proactive vs Reactive	<ul style="list-style-type: none"><li>• Reactive → Proactive → Predictive</li><li>• Monitoring – Infra, System &amp; Application</li><li>• Logging</li></ul>
Automation vs Manual	<ul style="list-style-type: none"><li>• Automate for efficiency and less human involvement (human error)</li></ul>
Reuse & Modularity	<ul style="list-style-type: none"><li>• Design for multiple use cases</li><li>• Adopt open-source</li><li>• Build modules and not a monolith</li></ul>
Simple over complex solutions	<ul style="list-style-type: none"><li>• Always opt for straightforward approach</li><li>• Reduce “moving parts” when possible</li></ul>

# Foundation of WB Architecture & DevOps Team

Mike Margolis | Chief Architect



Eitan Sela | Systems Architect



Reuven Kaplan | Systems Architect



Elran Regev | DevOps / DBA



July 2017

Stas Spivak | Systems Architect



Tal Klinger | IoT Architect



Leonid Kopylov | IT & Devops



# WINTER IS COMING



# Security

- Added additional auditing in AWS infrastructure (Cloudtrail)
- Removed users that don't absolutely need access and removing unused keys (AWS & DB)
- Moved to strict AWS IAM user/role management over systems.
- Integrated with SSO provider (OneLogin)
- Moved production environments to separate, isolated AWS account from Dev, Stage environments, and established additional DR account.
- Added additional Security Penetration testing processes to our software development practices
- Require full VPN access (highly restricted) to access any VPC resources.
- All AWS credentials in developers PC are obtained with OneLogin/STS, and expire after 12 hours.
- In later stages – GuardDuty, WAF, AWS Config, and much more...



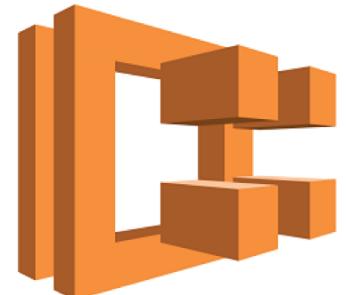
## CI/CD

- We provisioned Jenkins server to trigger automatic builds in Dev, manual in Prod.
- Java jars were now copied to the server and service restarted, without manually developers doing so.
- In later stages – Build and deploy containerized applications to ECS.
- In later stages – Use Jenkins jobs to enable triggering Lambda functions or long running ETLs by “clicking a button”.
- **Today we have more than 150 Jenkins jobs.**



# Containerization and micro services

- We started by splitting pos-engine, monolithic app, to three Spring Boot micro-services and deploy them to ECS.
- Move Analytics (data science) code in Python to be containerized.
  - All Linux cron jobs were converted to ECS Scheduled Tasks.
- In the following months, dozens of containerized micro-services will be developed and deployed to ECS.
- Use Fargate - For running POC Monthly reports once a month.
- **Today we have more than 50 Tasks Definitions and more than 100 containers deployed & running in production.**



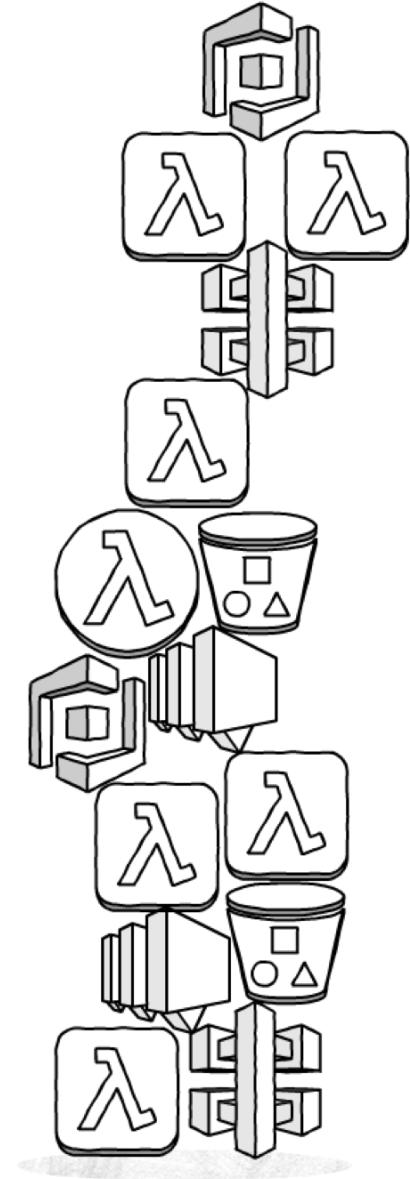
## Redshift as Managed DWH and development of ETLs

- We started using Redshift for our DWH:
  - Optimized for Business Intelligence (BI).
  - Columnar Storage for heavy analytics and aggregative queries.
- For our ETLs from MySQL to Redshift – Used PDI (Pentaho data integration) deployed as containerized Spring Boot app.
- **Result: Improved Dashboard and Mobile App query performance by several fold, and make improved user experience!**

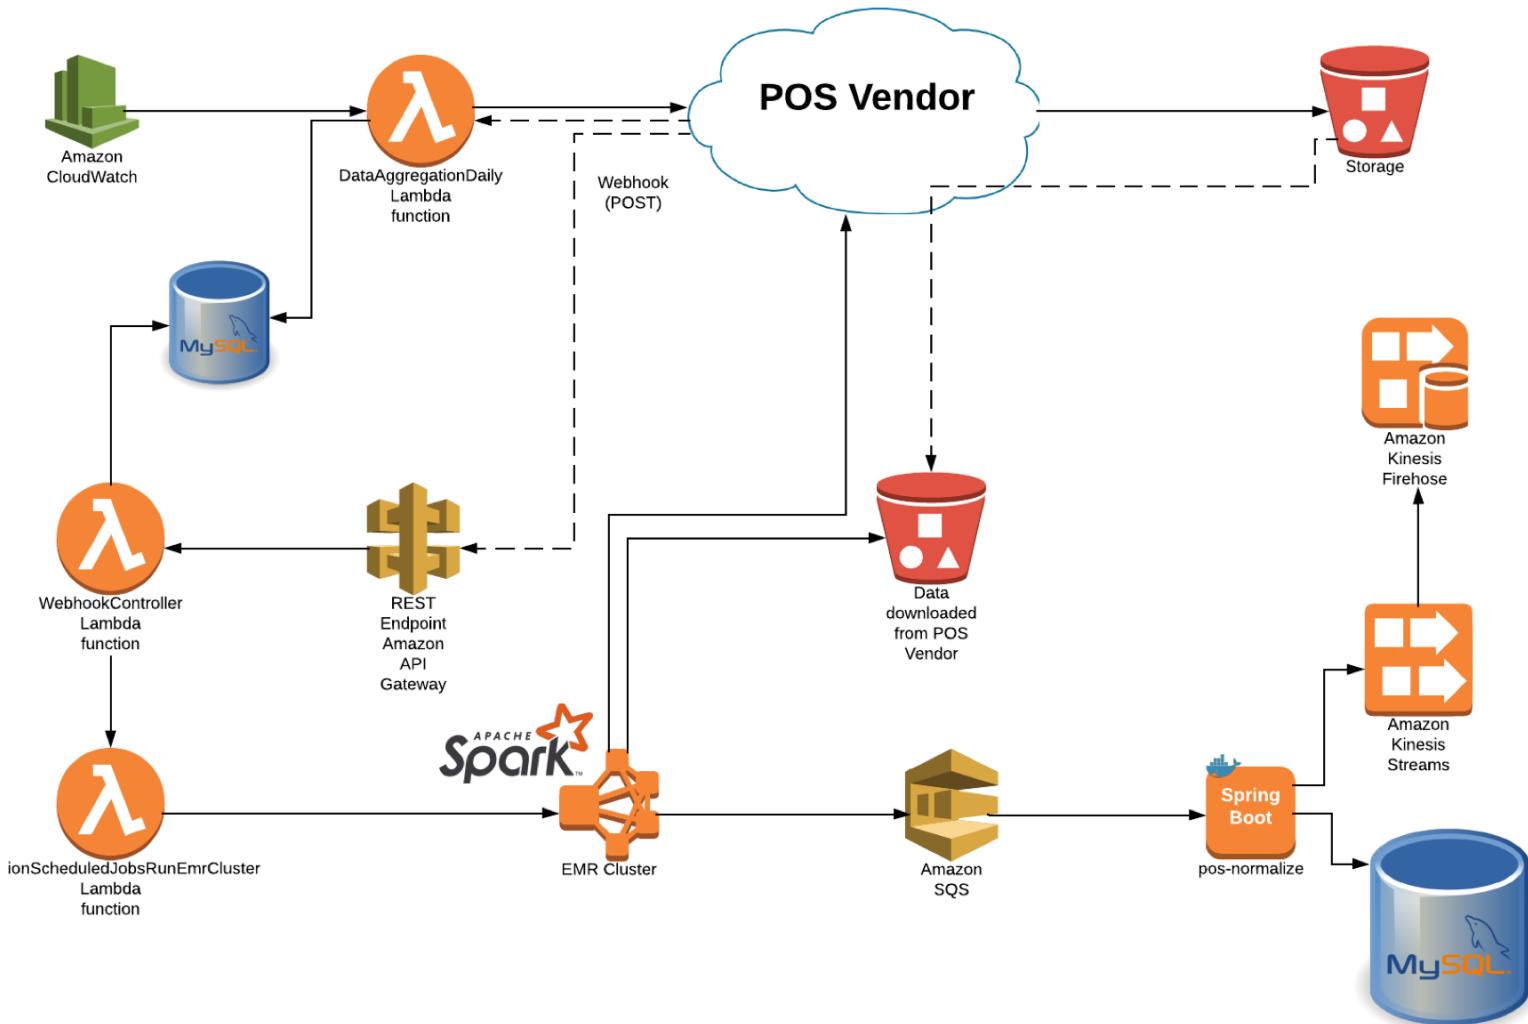


## Serverless - Start small, grow tall

- We use Lambda Functions extensively.
- Mostly with the CloudWatch Rule as trigger (cron like), but also with other triggers – S3, SQS, Kinesis, API Gateway, Lambda calling other Lambda and more.
- For Lambda CI/CD we use serverless framework.
- Lambdas are very easy to deploy – Just create *serverless.yml* and Jenkins job with few lines of shell script.
- For Lambda tracing, we use AWS X-Ray.
- **Today we have more than 90 Lambda Functions in production.**

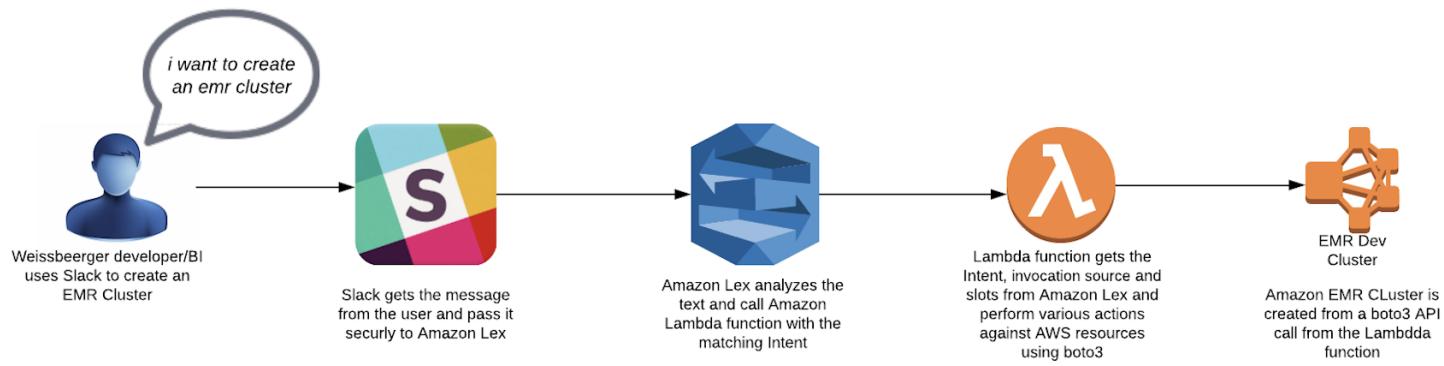


# Batch processing pipeline using Lambda Functions, CloudWatch and API Gateway



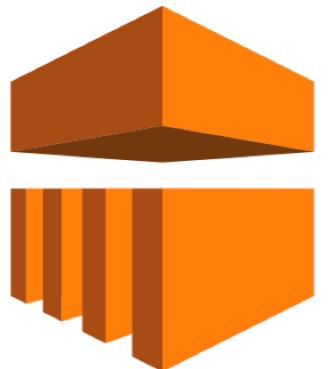
# AI & Machine learning

- From our home made algorithms, trained and deployed on EC2, and stored on MongoDB, to fully managed machine learning service using Amazon SageMaker.
- SageMaker Batch Transform Feature was developed as direct request from WB.
- We save 120 man-hours a month, provisioning EMR Clusters, using our “aws chatbot” built using Slack, Amazon Lex and Lambda functions.



# Implementing Big Data pipeline with EMR

- This effort started only on February 2018.
- There were two problems:
  - ETLs were fetching data directly from MySQL, which caused 100% CPU, slave lags and eventually, scaling up more and more, to enable normal work, thus exceeding budget.
  - BI devs had to wait sometimes hours, for heavy queries on MySQL to complete.
- After implementing full pipeline with Spark jobs ELTs that fetch data from Hive on S3:
  - ETLs (now ELTs) take minutes to complete. No stress on MySQL.
  - BI query on un-imagined amount of data, and wait usually less than 15 minutes.



## Implementing Big Data pipeline with EMR – con't

- We use transient EMR cluster and S3 as our Data Lake – Decoupling of storage from compute and data processing.
- We trained our BI to develop PySpark jobs. Our POS Java developers felt more comfortable developing Spark jobs in Scala, so we have also Scala.
- **Today we have more than 100 PySpark/Scala jobs and more than 40 developers, BI and Data engineers that work with the set of applications that are shipped with EMR.**



# EMR Clusters provisioned via WeissBeerger "aws chatbot" - A typical day

	Name	ID	Status	Creation time (UTC+2)	Elapsed time
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-pavel-cluster	j-2JI5CZ25KFBH1	Starting	2018-12-19 16:02 (UTC+2)	52 seconds
<input type="checkbox"/>	▶ <span style="color: green;">●</span> scheduled-jobs-run-emr-stg-hourly-pos-iot-all-jobs	j-23SVXQ4UYMMZZ	Running	2018-12-19 14:50 (UTC+2)	1 hour, 13 minutes
<input type="checkbox"/>	▶ <span style="color: orange;">●</span> on-demand-dev-emr-pavel-cluster	j-3UJMT455MW2B9	Terminating User request	2018-12-19 14:43 (UTC+2)	1 hour, 20 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-matanbiron-cluster	j-20OPRRAINRPAJ	Waiting Cluster ready	2018-12-19 13:35 (UTC+2)	2 hours, 28 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-lekha.b-cluster	j-16NYBI3D11VFK	Waiting Cluster ready	2018-12-19 13:32 (UTC+2)	2 hours, 30 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-avishai.ophel-cluster	j-23AC3931M38NI	Waiting Cluster ready	2018-12-19 13:21 (UTC+2)	2 hours, 42 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-matan.damary-cluster	j-26HBPO30ASV8J	Waiting Cluster ready	2018-12-19 10:51 (UTC+2)	5 hours, 11 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-hive-prod-replica-dev-emr-assaf.kazakov-cluster	j-MKECDOMZ8KK	Waiting Cluster ready	2018-12-19 10:17 (UTC+2)	5 hours, 45 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-hive-prod-replica-dev-emr-tal.barak-cluster	j-14F97HPY6KUFQ	Waiting Cluster ready	2018-12-19 10:15 (UTC+2)	5 hours, 48 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-hive-prod-replica-dev-emr-gal.geva-cluster	j-3TBMON7B16BGV	Waiting Cluster ready	2018-12-19 09:56 (UTC+2)	6 hours, 6 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-adi.shemesh-cluster	j-2271AWF4O1L27	Waiting Cluster ready	2018-12-19 09:41 (UTC+2)	6 hours, 22 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-michal-cluster	j-114LNW1IJQSEE	Waiting Cluster ready	2018-12-19 09:11 (UTC+2)	6 hours, 52 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-hadar_oren-cluster	j-MOBWF2F7CN6R	Waiting Cluster ready	2018-12-19 08:50 (UTC+2)	7 hours, 13 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-amit-cluster	j-32GNFMV4GI8HW	Waiting Cluster ready	2018-12-19 08:32 (UTC+2)	7 hours, 31 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-nevo.ezov-cluster	j-I53QQF2MZASJ	Waiting Cluster ready	2018-12-19 07:52 (UTC+2)	8 hours, 11 minutes
<input type="checkbox"/>	▶ <span style="color: green;">●</span> on-demand-dev-emr-eitan.sela-cluster	j-21PFRR0QPL5D0	Waiting Cluster ready	2018-12-19 07:26 (UTC+2)	8 hours, 37 minutes
...					

# Monitoring with Datadog

- All our monitoring is done in Datadog.
- We monitor:
  - Resource metrics - Physical components as CPU, memory, disks, and network interfaces.
  - Work metrics – Top-level health of the system - Throughput, Success/Error and Performance.
- Monitoring our containers deployed in ECS is done with dd-agent as side-car.

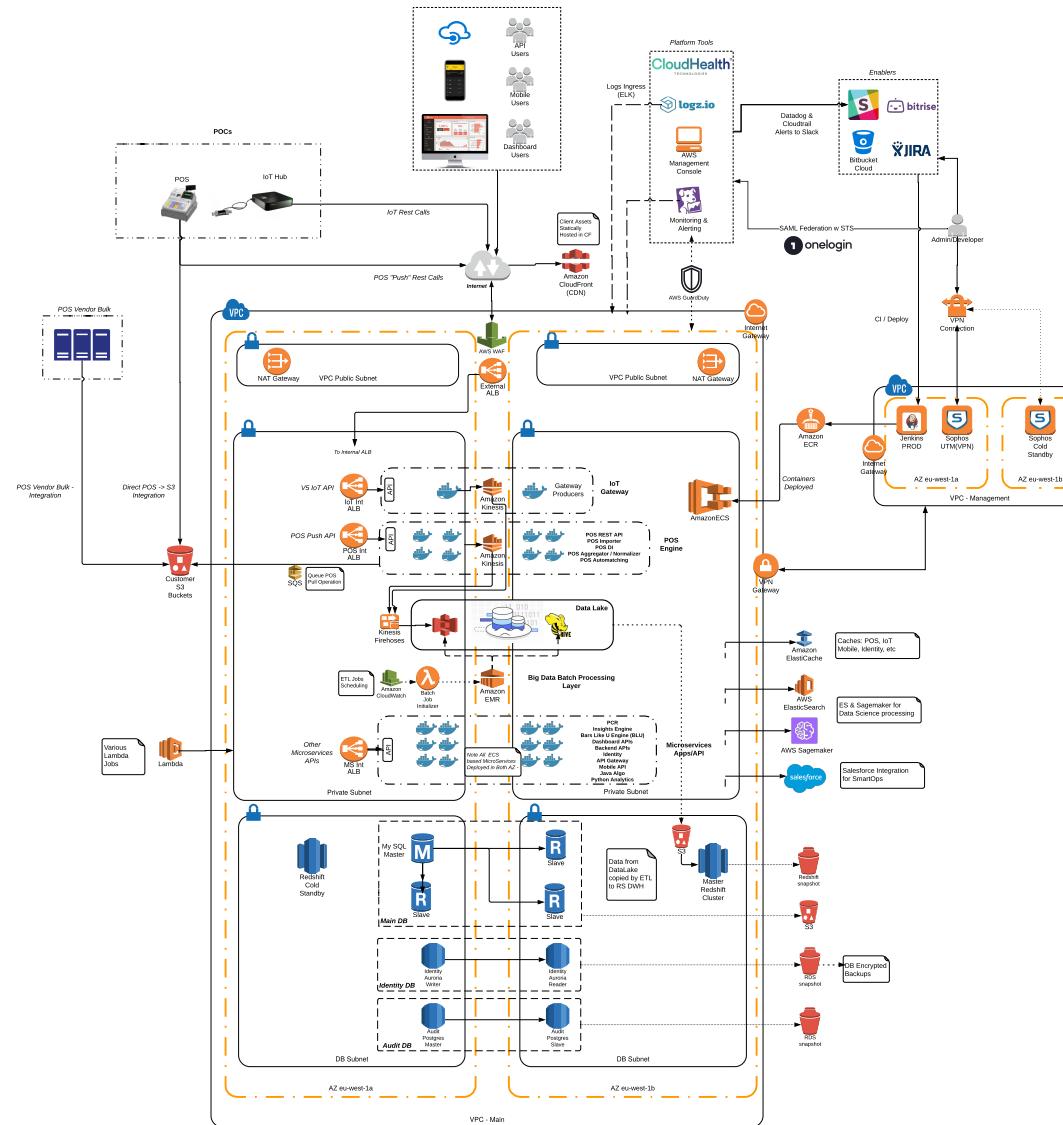


# Monitoring Everything – From heap size of Java apps to failed EMR steps

The screenshot shows the Datadog interface with the following details:

- Header:** Triggered Monitors, Manage Monitors (selected), Manage Downtime.
- Search Bar:** production status:ok
- Left Sidebar:** Watchdog, Events, Dashboards, Infrastructure, Monitors (selected), Metrics, Integrations, APM, Notebooks, Logs, Help, Team, eitan.sela.
- Monitors List:** Showing 50 of 63 results.
- Filters (Left):**
  - Status: Triggered (0), OK (63, checked), Muted (4).
  - Type: Integration (44), Event (6), Anomaly (5), Metric (3), Network (4), APM (0), Composite (0), Custom (0), Host (0), Process (1), Forecast (0), Live Process (0), Outlier (0).
  - Creator: Eitan Sela (37).
  - Tag: Filter tag list, \*.
- Table Headers:** STATUS, NAME ↑, TAGS.
- Table Data:** A list of 63 monitor entries, all marked as "OK". Some entries include icons like a red alert, a yellow warning, or a muted sound.

# WeissBeerger Architecture - December 2018



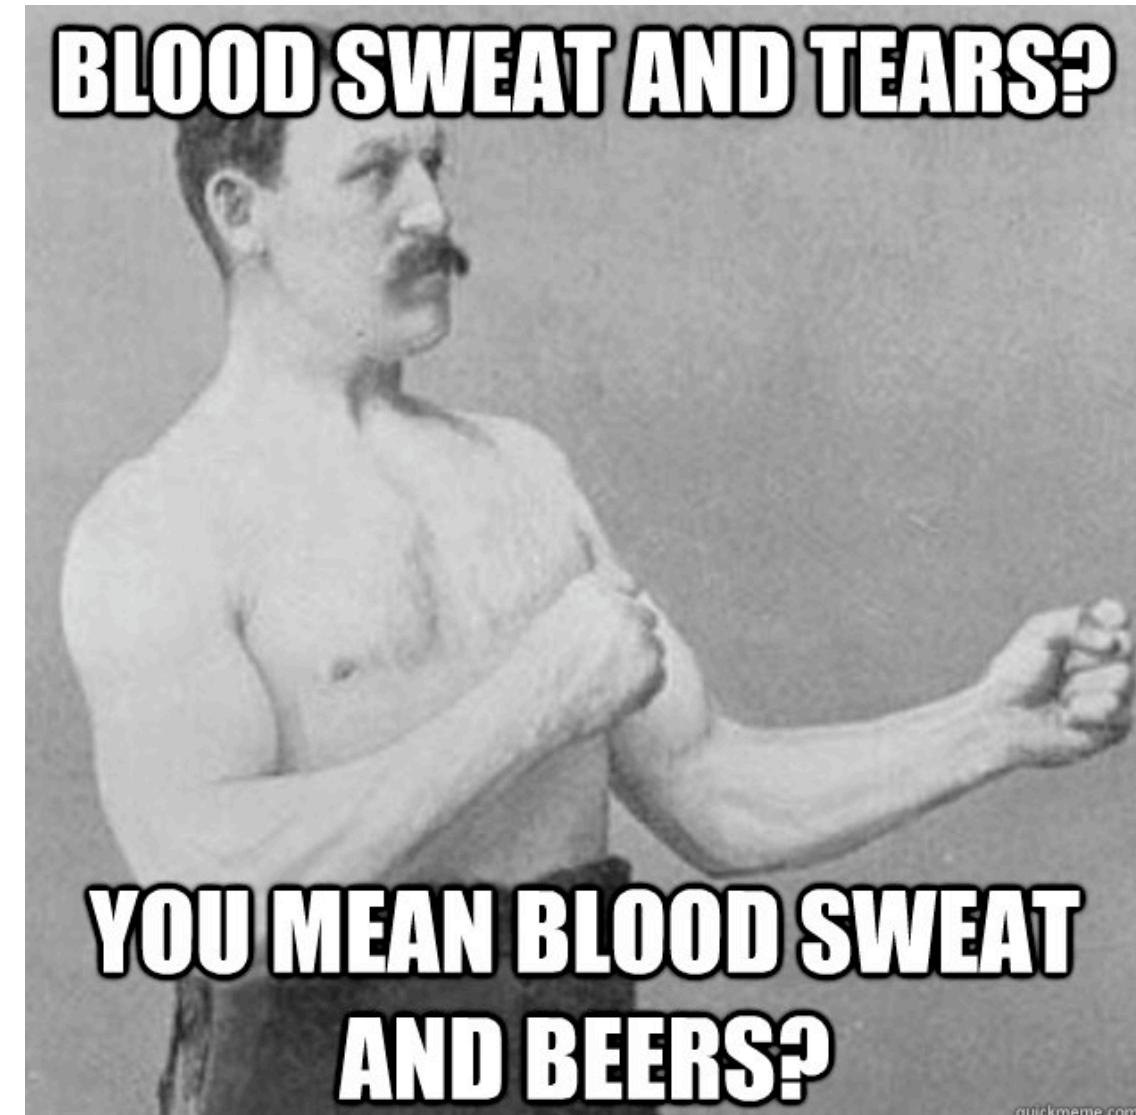
# AWS Usage – December 2018

- EC2
- Route 53
- ECS
- Fargate
- RDS
- Aurora
- SQS
- SNS
- Lambda
- CloudWatch
- CloudFront
- EMR
- Kinesis
- IoT Button
- SageMaker
- S3
- Glacier
- API Gateway
- GuardDuty
- WAF
- ALB
- Redshift
- DMS
- ElatstiCache
- Config
- X-Ray
- Lex
- WorkSpaces
- DynamoDB
- Textract – Future
- Timestream – Future
- EKS - Future



## How did we achieved all that?

- Commitment and hard work of architecture team.
- We were given mandate to do whatever we thought will help the company, considering costs and budget limits.
- Rapid adaptation of new technologies amongst R&D, DS and BI teams.



# The outcome – December 2017

## יצרנית הבירה הגדולה בעולם רכשה את ויסבירגר מט"

חברת הסטארט-אפ נמכרה בסכום מוערך של 80 מיליון דולר ■ ויסבירגר מספקת מערכות אונלייניקס מבוססות אינטרנט של הדברים לצרני משקאות, מסעדות וברים

א + א - 🔍

נתן יפתח, 26/01/2018 10:08



חחרות נשיאות בירה למראק בגרמניה / צילום: Michael Dalder, ריטרס

חברת הסטארט-אפ ויסבירגר מטל-אביב, המספקת מערכות אונלייניקס מבוססות אינטרנט של הדברים לצרני משקאות, מסעדות וברים, נמכרה בסכום מוערך של 80 מיליון דולר.

הרכשת היא אנհאוזר-בוש אינביו (Anheuser-Busch InBev), יצרנית הבירה הגדולה בעולם, המחזיקה כ-400 מותגי בירה, בהם בקס, באדוויזר, סטלה, הוגרדן וכו'.



## Industry with Heritage

**ABInBev** 600 years

 169 years

 152 years

## Global beer consumption

**187,000,000,000**  
**Liters consumed per year**

**50,000,000**  
**Outlets serving beverage**

## Future Challenges

- Move to fully IaC - Infrastructure as a code.
- Implement development & ops oriented culture change / Train our developers to gain further knowledge of AWS and how to leverage it.
- Improvements and costs savings on various Big Data services we use – implement AirFlow and start using Athena, automated cost analysis & remediation.
- Develop more for serverless.
- Move all our ML algorithms to SageMaker and develop new algos based on SageMaker pre made ones.
- Development of new self installed IoT hubs that uses AWS IoT and Greengrass.

# We Are Hiring!

- Senior Big Data Engineer.
- Senior Full Stack Developer.
- Data QA Engineer.
- Data Management Specialist.
- Head Of Commercial Analytics.
- Data Integrity Team Leader.
- Senior Product Manager.

[eitan.sela@weissbeerger.com](mailto:eitan.sela@weissbeerger.com)

<https://beverageanalytics.io/about-us/>



# Q & A